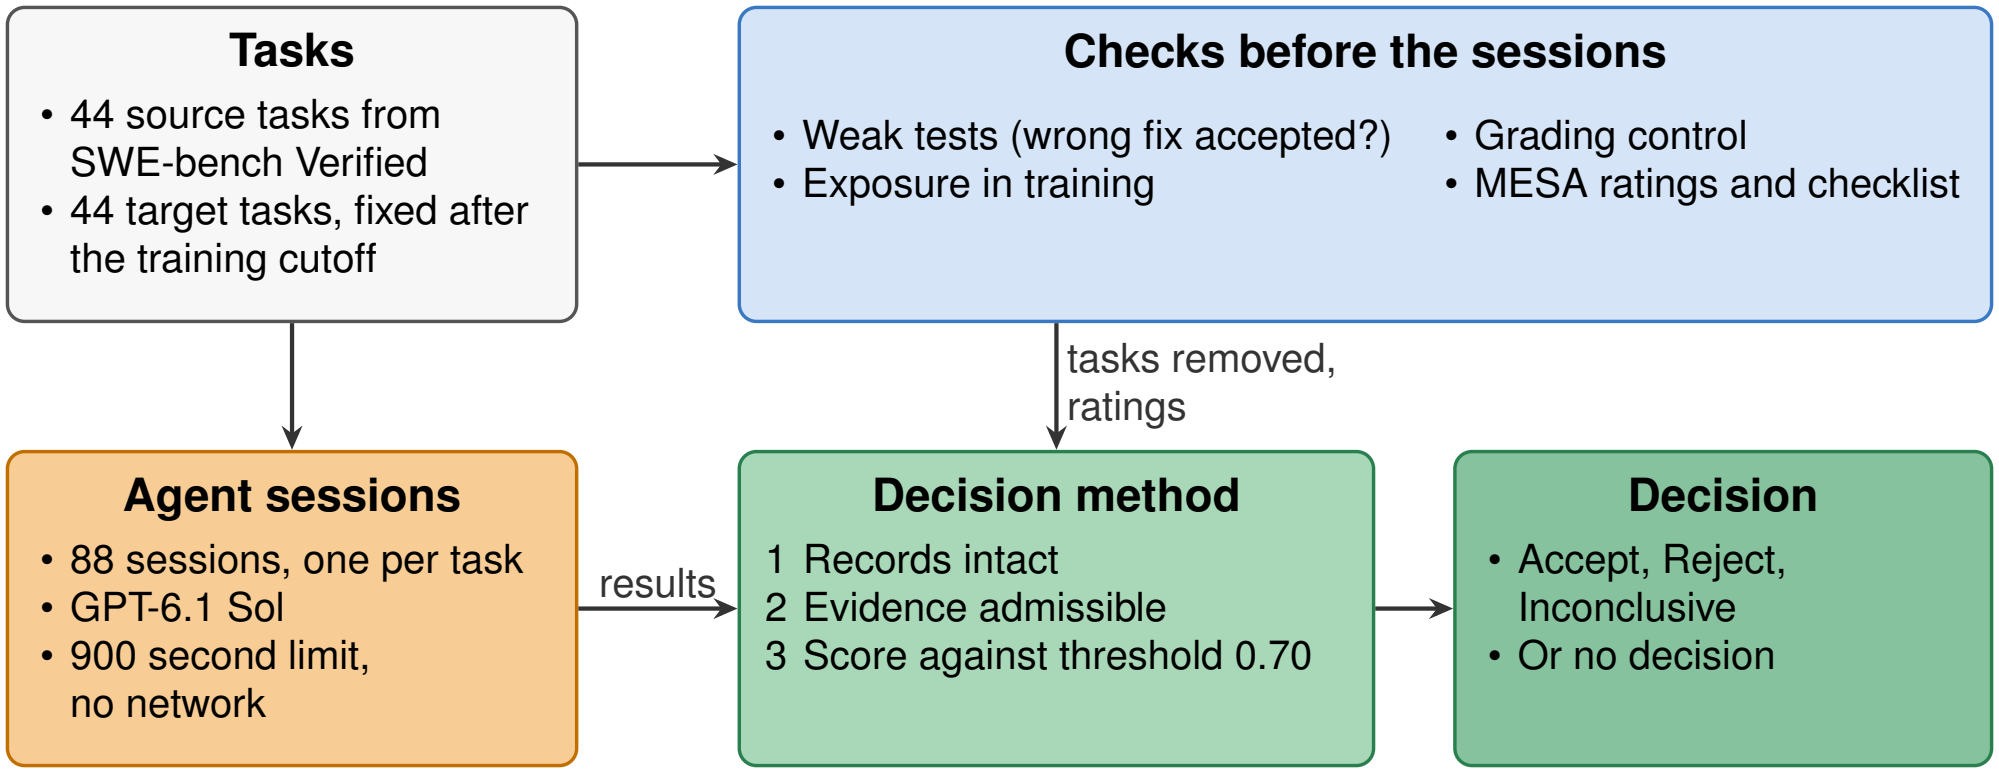
Tasks
• 44 source tasks from SWE-bench Verified
• 44 target tasks, fixed after the training cutoff
Checks before the sessions
• Weak tests (wrong fix accepted?)
• Exposure in training
• Grading control
• MESA ratings and checklist
tasks removed, ratings
Agent sessions
• 88 sessions, one per task
• GPT-6.1 Sol
• 900 second limit, no network
results
Decision method
1 Records intact
2 Evidence admissible
3 Score against threshold 0.70
Decision
• Accept, Reject, Inconclusive
• Or no decision